

# Advanced Enterprise Computing - Lecturenotes SoSe2016

Julius Hülsmann

20. Juli 2016

## Inhaltsverzeichnis

<b>1</b>	<b>Repetition</b>	<b>2</b>
1.1	ACID . . . . .	2
1.2	CAP . . . . .	2
1.3	PACELC . . . . .	2
1.4	Consistency . . . . .	3
1.4.1	Ordering and Data-centric consistency . . . . .	3
1.4.2	Ordering and Client-centric consistency . . . . .	3
<b>2</b>	<b>Replication and State Management</b>	
	<b>(25.04. - 09.05.)</b>	<b>4</b>
2.1	Motivation and Background . . . . .	4
2.1.1	Replication . . . . .	4
2.2	managing Replication . . . . .	6
2.2.1	Reconciliation . . . . .	8
2.2.2	Conclusion . . . . .	9
2.3	Paxos . . . . .	10
2.3.1	extensions . . . . .	10
2.3.2	Summary Paxos . . . . .	10
<b>3</b>	<b>Lecturenotes</b>	<b>11</b>
<b>4</b>	<b>Begriffe und Abkürzungen</b>	<b>11</b>

# 1 Repetition

## 1.1 ACID

**Atomitcity:** Either entire Transaction executed or aborted. Update all Replicas or none

**Consistency:** does not mean Data-Consistency but that the transaction produces consistent changes, client-centric / Data-centric, read-my-writes etc..

**Isolation:** Transactions are isolated from one another

**Durability:** Once the transaction is ready (commits) it remains.

Both the Atomitcity and the Isolation are managed by the **Transaction Manager**

- Acquires locks on behalf of the transactions
- guarantees serializable execution (= strongest form of isolation; outcome = outcome in case sequentially executed).  
guaranteed by use of 2PL

name	protocol's name	bsp implementations	distrib. syst
Atomacity	atomic commitment prot.	2PC	easy, expensive
Isolation	concurrency control protocol	2PL, Snapshot isolation	problematic

## 1.2 CAP

CAP = Consistency vs. Avaibility in case of Partitioning (Replication) Either one has to choose consistency or availability.

Gives information on the system's behavior in case of a system error.

## 1.3 PACELC

PACELC: partitioning: Avaibility/Consistency else Latency/Consistency

In case an update request arrives and the data is replicated and the system is working properly, there is a time difference between the moment the first replica receives the update and the other replicas are informed. Now it's the question whether to

1. immeadiately commit the new data (*chose the least latency*) or
2. whether not to respond until the 2nd replica respons (*chose consistency*)

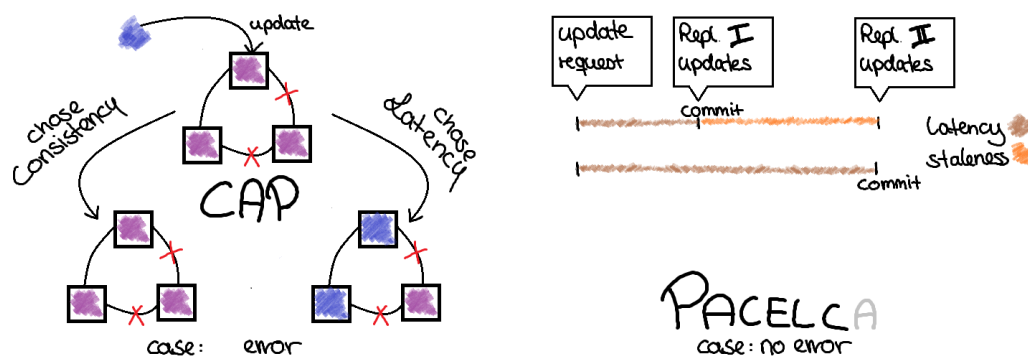


Abbildung 1: Cap and PACELC

## 1.4 Consistency

2 Dimensions:

- **Staleness** (how much is a given replica lagging behind) vs. **Ordering** (how much does the operation serializable order deviate among replicas)
- **Data-centric** and **Client-centric** consistency

### 1.4.1 Ordering and Data-centric consistency

- Sequential Consistency
- Causal consistency
- Eventual Consistency

### 1.4.2 Ordering and Client-centric consistency

- Monotonic Reads (never receive older values than previously read)
- Read-my-writes (not older than previously written)
- Write follows reads (only update Replicas that have at least got read entity)
- Monotonic Writes (update of one client: always executed sorted by time)

## 2 Replication and State Management (25.04. - 09.05.)

### 2.1 Motivation and Background

#### 2.1.1 Replication

**Definition - Replication** Process of maintaining multiple Copies of an Entity (Data / Process / File ...)

#### Advantages of Replication in General

- *System Availability / Fault tolerance / Security* in case
  - A Server fails
  - B Data is corrupted (vote against data, Byzantine)
- *Performance / Scalability*
  - A Workloads are spread across distributed Replicas
  - B Geodistribution for processing demands in client's proximity

#### Disadvantages of Replication in General

- Consistency vs. Performance

**Kinds of Replication** There are the following three different decisions one needs to take for developing a suitable Replication design:

1. „Physical Replication“(A), (B) or (C)
2. Defined access and update Mechanisms(Synchronous/Asynchronous PrimaryCopy/Update Everywhere)
3. Where to put the Replicas (Geo-distributed?)

**decision 1) Replication** In general there are the following kinds of „physical“ Replication. We do only consider (B).

#### decision 2) Replication Strategies

When	Sync./eager vs. Async./lazy	update propagation
Where	Primary Copy/master vs. upd. Everywhere/group	- location

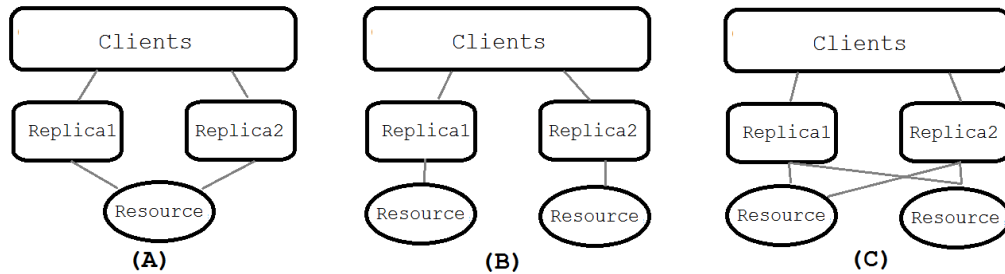


Abbildung 2: (A) Improves the availability only in case the Server Replicas use a replica cache coherence mechanism (B) Improves the availability. Usually ment by the term „Replication“.

**decision 3) Replicas' location** There are permanent Replicas, they can be geo-distributed or within the LAN  
Server-initiated (e.g. Push-Cache) or Client-initiated (e.g. Cache) temporary Replicas exist, too.

Push-Cache: Server record request-history

⇒ Tradeoff between consistency and latency is accentuated

**What happens to ACID in case of Replication?** Atomicity can be guaranteed using 2PC (but expensive) Problem: Serialization order must be the same at all replicas.

## 2.2 managing Replication

<p><u>Synchronous / eager (propagate before commit)</u></p> <ul style="list-style-type: none"> <li>Update coordination (ask every replica → overhead)</li> <li>+ Atomic Changes</li> <li>+ Strong Consistency</li> <li>- Availability</li> <li>- Latency</li> <li>- Execution + Response Time</li> <li>- complete ACID</li> </ul> <p><u>Asynchronous / lazy (commit immediately)</u></p> <ul style="list-style-type: none"> <li>- Not guaranteed</li> <li>- eventual consistency</li> <li>+ high availability</li> <li>+ low latency</li> <li>- Replication not transparent</li> </ul> <p><u>Primary Copy</u></p> <ul style="list-style-type: none"> <li>all secondary: read-only</li> <li>ACID can be guaranteed</li> <li>- low availability</li> <li>- Single Point of Failure</li> <li>- latency</li> <li>- bottleneck</li> </ul> <p><u>Update Everywhere</u></p> <ul style="list-style-type: none"> <li>Not guaranteed in case of simultaneous updates</li> <li>+ high avail.</li> <li>+ load evenly distributed</li> </ul>	Consistency	<p><u>Primary Copy</u></p> <ul style="list-style-type: none"> <li>+ globally consistent - expensive</li> <li>+ remote writes</li> <li>++ ACID</li> <li>- no coord.</li> <li>- longest response time</li> <li>- read only copies</li> <li>- low availability</li> <li>- does not scale</li> </ul>	<p><u>Update Everywhere</u></p> <ul style="list-style-type: none"> <li>- does not scale</li> <li>+ Isolation (each site as if it was only one, serializable)</li> <li>+ Consistency (etc.) can be achieved</li> <li>- coordination of updates</li> <li>- long response time</li> <li>- low availability</li> <li>- Symmetrical, elegant</li> <li>- does not scale</li> <li>- deadlocks</li> </ul>
		<p><u>Asynchronous</u></p> <ul style="list-style-type: none"> <li>- Inconsistency, Stale data</li> <li>- But no coord. of updates</li> <li>+ Short Response times</li> <li>- low write availability</li> <li>- inconsistent reads</li> <li>- limited scalability</li> </ul>	<ul style="list-style-type: none"> <li>- updates can be lost, conflicts</li> <li>- Incons.</li> <li>- coord. of updates</li> <li>+ Short Response time</li> <li>+ highest availability</li> <li>- inconsistent reads</li> <li>- reconciliation</li> <li>+ feasible in many applications (quick response)</li> </ul>

Abbildung 3: Different strategies: Advantages, Disadvantages, Properties

	<p><u>Primary Copy</u></p>	<p><u>Update Everywhere</u></p>
<u>Synchronous</u>	<p>Not used</p>	<ul style="list-style-type: none"> <li>• read one (locally) / write all</li> <li>• Isolation each site: 2PL</li> <li>• dealin' with site failures</li> <li>• site recovering</li> <li>• response time: msg overhead too high</li> <li>• transaction response time high</li> <li>• update 2N msg</li> <li>• Deadlocks (depend on <math>\frac{tx}{sec}</math>, db size, nr nodes)</li> <li>↳ System does not scale</li> </ul>
<u>Asynchronous</u>	<ul style="list-style-type: none"> <li>• Read locally</li> <li>• Prim. copy: local use of 2PL</li> </ul>	<ul style="list-style-type: none"> <li>• Read / Write, Transaction locally</li> <li>• local use of 2PL</li> <li>- no atomic commit prob but need 2 coordinate</li> <li>often words: Broadcast</li> <li>A B C</li> <li>update → X ← update</li> </ul>

Abbildung 4: Different strategies: Algorithm in detail

**Synchronous Update Everywhere Protocol** Assumption: All sites (Replicas) Contain the same data. Behavior if Transaction is to be executed

- Local use of 2PL for the following steps:
- READ only one site, in case the reading fails (timeout), read another copy
- WRITE at all sites (distributed locking protocol). This means that all copies of the data item need to lock the item (REQUEST, OBTAIN LOCK, ACK)

IF one site rejects, ABORT.

ADD All site not responding to a list of *missing writes*.

- VALIDATE (=commit) the transaction at the end, this means

IF NOT ALL the servers *missing writes* are down: ABORT

IF NOT ALL the servers that *accepted* are still available: ABORT

OTHERWISE Commits

⇒ Guarantees behaviour like if the sites were not replicated.

⇒ Execution is serializable

⇒ all Reads access the latest version

### Extensions for coping with failures

1. Site failure (reduces the availability)
  2. Behavior after the Site that failed is online again (outdated data available)
- Optimization: Most ideas based on Quorums

**Quorums** kind of a middleground between synchronous and Asynchronous updates.

Reads contact more than one Replica

Write contacts a quorum of Replicas.

Rules:

Read at least 1 Replica that has received the latest update:  $R + W > N$ ,

The minimum amount of writes must be greater than half of the amount of replicas  $\frac{W}{2} > N$

Quorums that don't follow these rules are called **sloppy Quorums**. (Dynamo + Cassandra)

Used to trade off read and write latency

**Different views on the subject** The solution to Replication strategies in *Database-POV* and *Distributed Systems-POV* differ, but have converged. Distributed System

- Set of Services implemented by Server Processes, invoked by client processes
- Each server has got a local state
- Group of servers / group communication – helps to reduce complexity
- Group communication primitives: provide 1toMany communication. Example: Atomic Broadcast (ACAST), View-Synchronous Broadcast (VSCAST)

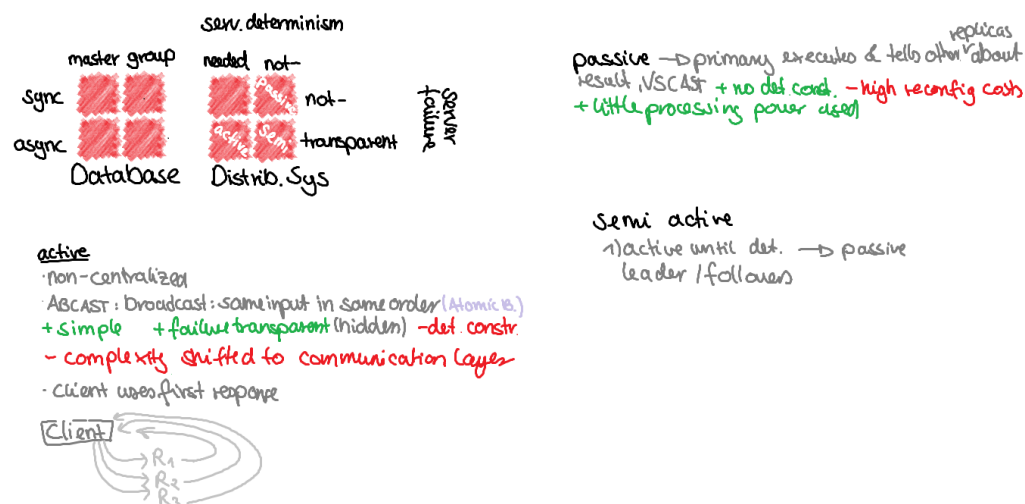


Abbildung 5: Distributed system perspective

### 2.2.1 Reconciliation

Who is responsible; When?

Alternative: Multiversion systems like DynamoDb

pre-arranged patterns:

- Last update wins (newer updates preferred over old ones)
- Site priority (preference to updates from headquarters)



- Largest value (the larger transaction is preferred)

ad-hoc decision making procedures

- Combination
- Analyze and remove not important ones
- ...

### **2.2.2 Conclusion**

if scalable  $\Rightarrow$  Conflicts have to be resolved / appear

## 2.3 Paxos

case: Conflicts cant be tolerated, using a scalable system. Not applicable for Byzantine failure

Rules:

1. exactly one value is chosen
2. non-triviality (chose one proposed val)
3. liveness (failure tolerant in case less than half fail)

Not nPC because may violate (Rules 1 and 3)

Client, Proposer, Acceptor, Learner

### 2.3.1 extensions

**Multi-Paxos**

**Fast Paxos**

**Generalized Paxos**

### 2.3.2 Summary Paxos

- ⇒ guarantees agreement in the presence of failures, Safety is always preserved
- ⇒ Conditions to affect liveness are hard to provoke
- ⇒ Uses same number of message rounds as 2PC

### 3 Lecturenotes

**Lecture 05?** start @ 81 Für Donnerstag paper mitbringen und Paxos anschauen.  
2016-05-09

**Paxos** (Represent as State-machine) - P. 77

Proposer

Phase 1 - Proposer choses Number largr than any value chosen before by Proposer. - Broadcast the integer *prepare(n)*, e.g. prepare(50)

Acceptors a) Not respond at all b) *reject* Reject, in case a higher value has been accepted. 50 ; something b) *prommise(n)* in case 50 ; everything. Also Send everything that has already been accepted.

If prposer receives majority of prommise resposns, -; proceed to Phase 2 ELSE -; Phase 1

Phase 2 - Check whether any ;n, value; have been returned. - YES: take max n's value - accept (n, value)

**Xtensions Paxos** **Multi-paxos** Determine Leader once Stay in phase 2, attatch the leader identifier Leader is the one to accept values

*Purpose: Optimize Speed* (get rid of the first phase, Master-Slave setup)

**Fast Paxos**

**Generalized Paxos** - Assumption: The execution order does not matter.

**CRDT** Conflict free / Communitive replicated Datatypes

Some operations are commutative, others not.

State- Based vs. Operation based.

theoretically it is possible to converge them but ... practice

IDEA INTEGER - example: e.g. not store int values but operations (increment / decrement))

SET - example

State - based Set

### 4 Begriffe und Abkürzungen

**Replication** Strategy to maintain mutiple copies of an entity on multiple Servers.

**Replica**

**CRDT** *conflict-free replicated data*

**Paxos**

**Commit** In case a Transaction commits, it is ready.

**Concurrency control protocol** guarantees isolation of Transactions

**2PL** Two phase locking (one concurrency control protocol)

**Snapshot Isolation** other concurrency control protocol implementation

**atomic commitment protocol** guarantees atomicity

**2PC** Two phase Commit

**Transaction Manager** Middleware Component; Manages Atomicity and Isolation of Transactions

**ACID** Atomicity + Consistency + Isolation + Durability

**serializability** a plan of executing multiple transactions in pseudo- parallel is called serializable in case the parallel execution comes to the same result as executing the transactions one after the other

**distributed locking protocol** 2PL z.B.? ??? Paxos??ß