Advanced Enterprise Computing - Lecturenotes SoSe2016

Julius Hülsmann

20. Juli 2016

Inhaltsverzeichnis

1   Repetition

1.1   ACID

Atomiticity: Either entire Transaction executed or aborted. Update all Replicas or none
Consistency: does not mean Data-Consistency but that the transaction produces consistent changes, client-centric / Data-centric, read-my-writes etc..
Isolation: Transactions are isolated from one another
Durability: Once the transaction is ready (commits) it remains.

Both the Atomitcity and the Isolation are managed by the **Transaction Manager**

- Aquires locks on behalf of the transactions

- guarantees serializable execution (= strongest form of isolation; outcome = outcome in case sequentially executed).
  guaranteed by use of 2PL

name	protocol's name	bsp implementations	distrib. syst
Atomacity	atomic committement prot.	2PC	easy, expensive
Isolation	concurrency control protocol	2PL, Snapshot isolation	problematic

1.2   CAP

CAP = Consistency vs. Avaiability in case of Partitioning (Replication) Either one has to choose consistency or availability.
Gives information on the system's behavior in case of a system error.

1.3   PACELC

PACELC: partitioning: Avaibaility/Consistency else Latency/Consistency
In case an update request arrives and the data is replicated and the system is working properly, there is a time difference between the moment the first replica receives the update and the other replicas are informed. Now it's the question whether to

1. immeadiately commit the new data *(chose the least latency)* or

2. whether not to respond until the 2nd replica respons *(chose consistency)*
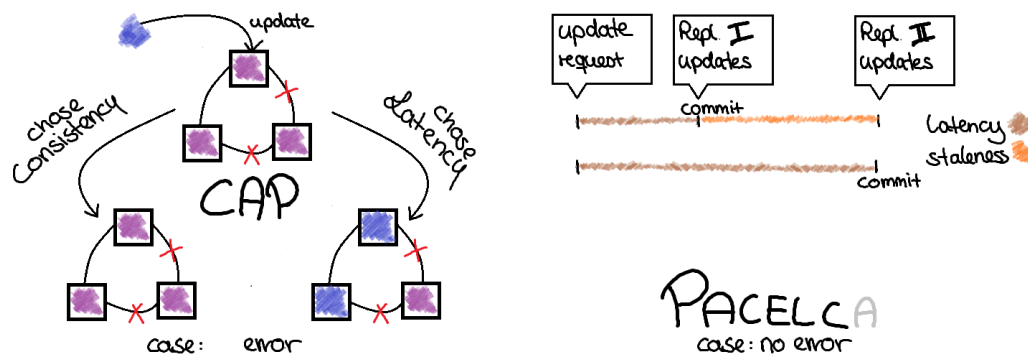
Abbildung 1: Cap and PACELC

1.4   Consistency

2 Dimensions:

- **Staleness** (how much is a given replica lagging behind) vs. **Ordering** (how much does the operation serializable order deviate among replicas)

- **Data-centric** and **Client-centric** consistency

1.4.1   Ordering and Data-centric consistency

- Sequential Consistency

- Causal consistency

- Eventual Consistency

1.4.2   Ordering and Client-centric consistency

- Monotonic Reads (never receive older values than previously read)

- Read-my-writes (not older than previously written)

- Write follows reads (only update Replicas that have at least got read entity)

- Monotonic Writes (update of one client: always executed sorted by time)

2   Replication and State Management (25.04. - 09.05.)

2.1   Motivation and Background

2.1.1   Replication

Definition - Replication   Process of maintaining multiple Copies of an Entity (Data / Process / File ...)

Advantages of Replication in General

- *System Availability / Fault tolerance / Security* in case
  A   Server fails
  B   Data is corrputed (vote against data, Byzantine)

- *Performance / Scalability*
  A   Workloads are spread across distributed Replicas
  B   Geodistribution for processing demands in client's proximity

Disadvantages of Replication in General   Consistency / Performance

Kinds of Replication   There are the following three different decisions one needs to take for developing a suitable Replication design:

1. „Phyiscal Replication"(A), (B) or (C)

2. Defined access and update Mechanisms(Synchronous/Asynchronous PrimaryCopy/Update Everywhere)

3. Where to put the Replicas (Geo-distributed?)

Abbildung 2: (A) Improves availability if Server Replicas use replica cache coherence mechanism (B)Improves availability. Meant by „Replication"4us.

decision 1) Replication   In general there are the following kinds of „physical" Replication. We do only consider (B).

decision 2) Replication Strategies
When   Sync./eager vs. Async./lazy   update propagation
Where   Primary Copy/master vs. upd. Everywhere/group   - location

decision 3) Replicas' location   There are permanent Replicas, they can be geo-distributed or within the LAN
Server-initiated (e.g. Push-Cache) or Client-initiated (e.g. Cache) temporary Replicas exist, too.
Push-Cache: Server record request-history
⇒ Tradeoff between consistency and latency is accentuated

What happens to ACID in case of Replication?   Atomicity can be guaranteed using 2PC (but expensive) Problem: Serialization order must be the same at all replicas.

2.2 managing Replication

Abbildung 3: Different strategies: Advantages, Disadvantages, Properties

Abbildung 4: Different strategies: Algorithm in detail

Synchronous Update Everywhere   Protocol Assumption: All sites (Replicas) Contain the same data. Behavior if Transaction is to be executed

- Local use of 2PL for the following steps:

- READ only one site, in case the reading fails (timeout), read another copy

- WRITE at all sites (distributed locking protocol). This means that all copies of the data item need to lock the item (REQUEST, OBTAIN LOCK, ACK)
  .
  IF one site rejects, ABORT.
  ADD All site not responding to a list of *missing writes*.

- VALIDATE (=commit) the transaction at the end, this means

  IF NOT ALL the servers *missing writes* are down: ABORT
  IF NOT ALL the servers that *accepted* are still available: ABORT
  OTHERWISE Commits

⇒ Guarantees behaviour like if the sites were not replicated.

⇒ Execution is serializable

⇒ all Reads access the latest version

Extensions for coping with failures
1. Site failure (reduces the avalability)
2. Behavior after the Site that failed is online again (outdated data available)
Optimization: Most ideas based on Quorums

Quorums   kind of a middleground between syncrhonous and Asynchronous updates.
Reads contact more than one Replica
Write contacts a quorum of Replicas.
Rules:
Read at least 1 Replica that has received the latest update: $R + W > N$,
The minimum amount of writes must be greater than half of the amount of replicas $\frac{W}{2} > N$
Quorums that don't follow these rules are called **sloppy Quroums**. (Dynamo + Cassandra)
Used to trade off read and write latency

Different views on the subject   The solution to Replication strategies in *Database-POV* and *Distributed Systems-POV* differ, but have converged. Distributed System

- Set of Services implemented by Server Processes, invoked by client processes

- Each server has got a local state

- Group of servers / group communication – helps to reduce complexity

- Group communciation primitives: provide 1toMany communication. Example: Atomic Broadcast (ACAST), View-Syncrhonous Broadcast (VSCAST)
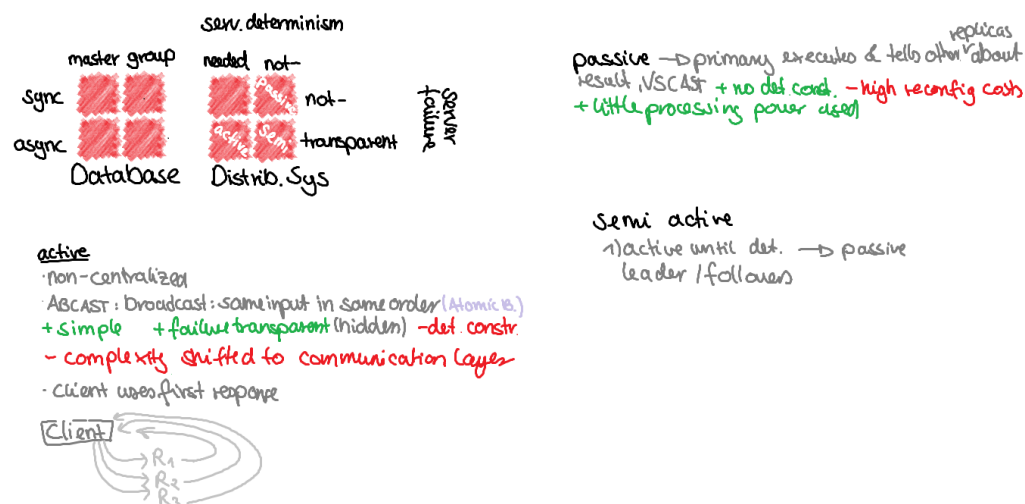
Abbildung 5: Distributed system perspective

2.2.1   Reconciliation

Who is responsible; When?
Alternative: Multiversion systems like DynamoDb
pre-arranged patterns:

- Last update wins (newer updates preferred over old ones)

- Site priority (preference to updates from headquarters)

- Largest value (the larger transaction is preferred

ad-hoc decision making procedures

- Combination

- Analyze and remove not important ones

- ...

2.2.2 Conclusion

if scalable $\Rightarrow$ Conflicts have to be resolved / appear

2.3 Paxos

case: Conflicts cant be tolerated, using a scalable system. Not applicable for Byzantine failure
Rules:

1. exactly one value is chosen

2. non-triviality (chose one proposed val)

3. liveness (failure tolerant in case less than half fail)

Not nPC because may violate (Rules 1 and 3)
Client, Proposer, Acceptor, Learner

Abbildung 6: arg2

2.3.1 extensions

Multi-Paxos Phases 1 + 2 to determine Leader
– Basically, a master-slave setup with fail-over strategy for the master

Fast Paxos Phases 1 + 2 to determine Leader
Clients directly issue accept requests
Leader reconsiles if necessary

Generalized Paxos Phases 1 and 2 determine leader
– Clients directly issue propose requests, acceptors send accepted requests to leader and learners
– Leader reconciles if necessary - Assume: The execution Order does not matter

2.3.2 Summary Paxos

⇒ guarantees agreement in the presence of failures, Safety is always preserved

⇒ Conditions to affect liveness are hard to provoke

⇒ Uses same number of message rounds as 2PC

2.4 Conflict free Replicated data types

Situation Async Everywhere, Staleness + Ordering, Ordering to be delt with
Idea: Design concurrent operations as non-conflicting

1. State-based
   Define a deterministic merge function that fulfills a set of formal requirements

2. Operation-based
   Define commutative operations that are called

both can theoretically converge

2.4.1 Summary CRDTs

Advantage - Ensure convergence - without Application level conflict resolution
- without Low performance due to consensus
- without Lost updates,
- Hide complexity of building distributed systems
- Allow clear design of concurrent semantics

Disadvantage - Limited in applicability due to - lack of global invariants
- supported data types and semantics,
- performance constraints (communication overhead)
- storage volume constraints (tombstones, internal type overhead)

Conclusion ⇒ An elegant solution for special use cases but cannot replace
general purpose database systems

3 Begriffe und Abkürzungen

Replication Strategy to maintain mutiple copies of an entity on multiple Servers.

Replica

CRDT *conflict-free replicated data*

Paxos

Commit In case a Transaction commits, it is ready.

Concurrency control protocol guarantees isolation of Transactions

2PL Two phase locking (one concurrency control protocol)

Snapshot Isolation other cuncurrency control protocol implementation

atomic commitement protocol guarantees atomaticity

2PC Two phase Commit

Transaction Manager Middleware Component; Manages Atomacity and Isolation of Transactions

ACID Atomacity + Consistency + Isolation + Durability

serializability a plan of executing multiple transactions in pseudo- parallel is called serializable in case the parallel execution comes to the same result as executing the transactions one after the other

distributed locking protocol 2PL z.b.? ??? Paxos??ß